

# Computational Bayesian data analysis

Stan language

---

Bruno Nicenboim / Shravan Vasishth

2020-03-11



Stan is (mostly) written in C++ and can be accessed through several interfaces:

- R
- Python
- Matlab
- Stata
- bash
- etc

# A Stan program

- usually saved as a `.stan` file
- accessed through R (or other interfaces)
- organized into a sequence of optional and obligatory blocks, which must be written in order.

# A Stan program

```
functions {  
  // This is an optional block used for functions that can be used in other blocks.  
}  
data {  
  // Obligatory block that specifies the required data for the model.  
}  
transformed data {  
  // Optional block if we want to manipulate the data.  
}  
parameters {  
  // Obligatory block that specifies the model's parameters.  
}  
transformed parameters {  
  // Optional block if we want to manipulate the parameters (re-parametrize the model).  
}  
model {  
  // Obligatory block that specifies the model's likelihood and priors.  
}  
generated quantities {  
  // Optional block if we want to manipulate the output of our model.  
}
```

## A Stan program

- every variable used needs to be declared first with its type (real, integer, vector, matrix, etc.).
- there must be a semi-colon (;) at the end of each line.

**Some examples of data type below (but check the [Stan reference manual](#) for more):**

- Variable `mu` contains a real number, either positive or negative:

```
real mu;
```

- Variable `X` contains a real number that is bounded between two numbers (or by only one). Suppose `X` is some type of measurement that can only be between 0 and 1000. (We can add lower and/or upper to any type.)

```
real<lower = 0, upper = 1000> X;
```

- Variable N contains integers, such as the number of observations (which should be  $> 0$ ):

```
int<lower = 0> N;
```

- For vectors, row\_vectors and matrices, we must define the number of elements (in each margin) that they will contain. This number can be defined earlier (such as N, the number of observations, in our example). (In Stan, the values inside a vectors and matrices are always *real*.)

```
vector<lower = 0> [N] Y;
```

```
row_vector<lower = 0> [10] Y;
```

```
matrix<upper = 0> [3, J] Rho;
```



- Any type can be converted into an array of as many dimensions as we want, even vectors and matrices. It's worthwhile to take a look at [Array Data Types section of the Stan reference manual](#).

```
real mu[2]; // one dimension, two places
```

```
int<lower = 0> N[x,y,z]; // 3 dimensions, with x, y, z places.
```

```
vector<lower = 0> [N] Y[2]; // array of one dimensions that  
contain 2 vectors of N places
```

## Example: Cloze probability with Stan (Binomial likelihood)

**We want to derive the posterior distribution of the Cloze probability of “umbrella”,  $\theta$ :**

- Data: a word (e.g., “umbrella”) was answered 80 out of 100 times,
- Likelihood: a binomial distribution
- Prior for  $\theta$ :  $Beta(a = 4, b = 4)$

## Example: Cloze probability with Stan (Binomial likelihood)

```
data {  
  int<lower = 1> N;  // Total number of answers  
  int<lower = 0, upper = N> k;  // Number of times "umbrella" was answered  
}  
parameters {  
  // theta is a probability, it has to be constrained between 0 and 1  
  real<lower = 0, upper = 1> theta;  
}  
model {  
  // Prior on theta:  
  target += beta_lpdf(theta | 4, 4);  
  // Likelihood:  
  target += binomial_lpmf(k | N, theta);  
}
```

(\*) every statement with `target +=` increments the unnormalized *log* posterior probability.

- Save the previous model as `stan_models/binomial.stan` (don't run it in R).
- Use the following code to call the model from R:


```
library(rstan)
options(mc.cores = parallel::detectCores())
lst_cloze_data <- list(k = 80, N = 100)
# Fit the model with the default values of number
# of chains and iterations (chains = 4, iter = 2000)
fit_cloze <- stan(
  file = "stan_models/binomial_cloze.stan",
  data = lst_cloze_data
)
```

```
fit_cloze
```

```
## Inference for Stan model: binomial_cloze.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##           mean se_mean   sd  2.5%  25%  50%  75%  98%
## theta  0.78     0.00 0.04  0.69  0.75  0.78  0.8  0.85
## lp__ -5.05     0.02 0.73 -7.15 -5.23 -4.77 -4.6 -4.54
##           n_eff Rhat
## theta  1412     1
## lp__   1922     1
##
## Samples were drawn using NUTS(diag_e) at Wed Mar 11 15:18:52 2020.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```

bayesplot (Gabry and Mahr 2019) is a wrapper around ggplot2 (Wickham et al. 2019) and has several convenient functions to plot the samples (see their [vignette](#)).

```
library(bayesplot)
# We need to convert the fit to plot it with bayesplot functions
df_fit_cloze <- as.data.frame(fit_cloze)
# Bayes plot functions start with mcmc_
mcmc_dens(df_fit_cloze, pars = "theta") +
  geom_vline(xintercept = mean(df_fit_cloze$theta))
```



## A more complex Stan model

$$\begin{aligned}rt_n &\sim \text{LogNormal}(\mu, \sigma) \\ \mu &\sim \text{Normal}(6, 1.5) \\ \sigma &\sim \text{Normal}_+(0, 1)\end{aligned}\tag{1}$$

# A more complex Stan model

```
data {  
  int<lower=1> N_obs;  
  vector[N_obs] rt;  
}  
parameters {  
  real<lower=0> sigma;  
  real mu;  
}  
model {  
  target += normal_lpdf(mu | 6, 1.5);  
  target += normal_lpdf(sigma | 0, 1) -  
    normal_lccdf(0 | 0, 1);  
  target += lognormal_lpdf(rt | mu, sigma);  
}  
generated quantities {  
  vector[N_obs] rt_sim;  
  // posterior predictive  
  for(i in 1:N_obs) {  
    rt_sim[i] = lognormal_rng(mu, sigma);  
  }  
}
```



```
library(dplyr)
library(readr)
df_noreading_data <-
  read_csv("./data/button_press.csv")
lst_noreading <- list(
  N_obs = nrow(df_noreading_data),
  rt = df_noreading_data$rt
)

fit_lognormal_reading <- stan("stan_models/lognormal.stan",
  data = lst_noreading
)
```

```
print(fit_lognormal_reading, pars = c("sigma", "mu"))
```

```
## Inference for Stan model: lognormal.
```

```
## 4 chains, each with iter=2000; warmup=1000; thin=1;
```

```
## post-warmup draws per chain=1000, total post-warmup draws=4000.
```

```
##
```

```
##      mean se_mean   sd 2.5%  25%  50%  75%  98% n_eff
```

```
## sigma 0.13      0 0.01 0.13 0.13 0.13 0.14 0.15  3081
```

```
## mu    5.12      0 0.01 5.10 5.11 5.12 5.12 5.13  4482
```

```
##      Rhat
```

```
## sigma    1
```

```
## mu       1
```

```
##
```

```
## Samples were drawn using NUTS(diag_e) at Wed Mar 11 15:46:49 2020.
```

```
## For each parameter, n_eff is a crude measure of effective sample size,
```

```
## and Rhat is the potential scale reduction factor on split chains (at
```

```
## convergence, Rhat=1).
```

# Extracting Stan code from brms

```
library(brms)
fit_press_ln <- brm(rt ~ 1,
  data = df_noreading_data,
  family = lognormal(),
  prior = c(
    prior(normal(6, 1.5), class = Intercept),
    prior(normal(0, 1), class = sigma)
  )
)
```

## Extracting code when the model was run:

```
stancode(fit_press_ln)
```

```

## // generated with brms 2.12.0
## functions {
## }
## data {
##   int<lower=1> N; // number of observations
##   vector[N] Y; // response variable
##   int prior_only; // should the likelihood be ignored?
## }
## transformed data {
## }
## parameters {
##   real Intercept; // temporary intercept for centered predictors
##   real<lower=0> sigma; // residual SD
## }
## transformed parameters {
## }
## model {

...

```

...

```
## }  
## model {  
##   // initialize linear predictor term  
##   vector[N] mu = Intercept + rep_vector(0, N);  
##   // priors including all constants  
##   target += normal_lpdf(Intercept | 6, 1.5);  
##   target += normal_lpdf(sigma | 0, 1)  
##     - 1 * normal_lccdf(0 | 0, 1);  
##   // likelihood including all constants  
##   if (!prior_only) {  
##     target += lognormal_lpdf(Y | mu, sigma);  
##   }  
## }  
## generated quantities {  
##   // actual population-level intercept  
##   real b_Intercept = Intercept;  
## }
```

## Extracting code before the model was run:

```
make_stancode(rt ~ 1,  
  data = df_noreading_data,  
  family = lognormal(),  
  prior = c(  
    prior(normal(6, 1.5), class = Intercept),  
    prior(normal(0, 1), class = sigma)  
  )  
)
```

## Extracting the data when the model was run:

```
ls_stan <- standata(fit_press_ln)
ls_stan %>% str()

## List of 5
## $ N          : int 361
## $ Y          : num [1:361(1d)] 141 138 128 132 126 134 163 149 133 110 ...
## $ K          : int 1
## $ X          : num [1:361, 1] 1 1 1 1 1 1 1 1 1 1 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:361] "1" "2" "3" "4" ...
## .. ..$ : chr "Intercept"
## ..- attr(*, "assign")= int 0
## $ prior_only: int 0
## - attr(*, "class")= chr "standata"
```



## Extracting the data before the model was run:

```
ls_stan <- make_standata(rt ~ 1,  
  data = df_noreading_data,  
  family = lognormal(),  
  prior = c(  
    prior(normal(6, 1.5), class = Intercept),  
    prior(normal(0, 1), class = sigma)  
  )  
)
```

# References

Gabry, Jonah, and Tristan Mahr. 2019. *Bayesplot: Plotting for Bayesian Models*. <https://CRAN.R-project.org/package=bayesplot>.

Wickham, Hadley, Winston Chang, Lionel Henry, Thomas Lin Pedersen, Kohske Takahashi, Claus Wilke, Kara Woo, and Hiroaki Yutani. 2019. *Ggplot2: Create Elegant Data Visualisations Using the Grammar of Graphics*. <https://CRAN.R-project.org/package=ggplot2>.